Ontology-Constrained Multi-Agent Simulation for Adversarial Opinion Susceptibility Auditing: Analytical Guardrails, Profile Feature Network Diagnostics, and Execution Integrity

Stijn Van Severen[1,*] Thomas De Schryver[1] Mira Ostyn[1]

[1] Department of Applied Mathematics, Computer Science and Statistics, Ghent University, Ghent, Belgium

[*] Correspondence: stijn.vanseveren@ugent.be

Keywords: cognitive security; influence operations; ontology engineering; multi-agent simulation; structural equation modeling; network analysis; adversarial effectivity

Abstract

This paper presents and validates an ontology-driven pipeline for susceptibility auditing under adversarial political persuasion. The design crossed 25 deterministic pseudoprofiles with four attack leaves and ten opinion leaves sampled across four policy domains, yielding 1,000 attacked scenarios. The profile ontology retained survey-mappable Big Five, socioeconomic status, political psychology, social context, and demographic features (91 modeled dimensions after exclusion of earlier noise-heavy inventories). The analysis stack combined repeated-outcome structural equation modeling, conditional susceptibility estimation via task-specific ridge models, hierarchical importance decomposition, bootstrap rank uncertainty quantification, and a profile feature network built from the deterministic profile data.

The central finding is methodological: during the present study, opinion-generation stages fell back to deterministic substitute outputs after API-credit exhaustion, producing 100% fallback use in baseline and post-exposure assessment and zero-valued realism scores. Under those degraded execution conditions, mean adversarial effectivity was 0.10, 49.6% of rows had positive adversarial effectivity, ICC(1) for $|\Delta|$ was 0.011, ridge cross-validated R^2 was -0.175, random-forest out-of-bag R^2 was -0.194, and the elastic net selected zero features. The robustness layer behaved appropriately: it surfaced weak-signal diagnostics instead of supporting unwarranted moderator claims.

One component remained fully interpretable despite the execution failure: the deterministic profile feature network. The resulting graph contained 91 nodes and 1,807 edges at $|\rho| \geq .15$, with density = .441, five communities, and modularity = .313. Online-behavior, institutional-trust, and Big Five openness/extraversion blocks formed the most connected regions of the feature space, while cross-community bridge diagnostics highlighted variables that link otherwise distinct ontology families. This study demonstrates that the analytical guardrails, visualization layer, and network diagnostics of the platform function correctly under degraded execution, and that interpretable structure is recoverable from deterministic components even when opinion-generation fails.

Introduction

Online influence operations exploit attention, trust, social proof, and identity-congruent framing rather than technical intrusion alone (Bennett and Livingston, 2018; Lazer et al., 2018; Lewandowsky et al., 2017). In that setting, a useful defensive question is not merely whether manipulative content exists, but whether specific attack mechanisms shift opinion in adversarially favorable directions for some profiles more than for others. Psychographic tailoring has made this question operationally salient because persuasive messages aligned to personality structure can outperform untailored variants (Matz et al., 2017). At the same time, resilience interventions increasingly focus on strengthening citizens' ability to recognize and resist manipulation rather than treating misinformation exposure as a purely platform-level problem (Kozyreva et al., 2020; Pennycook and Rand, 2021; Roozenbeek et al., 2022).

The present study addresses that defensive problem through an ontology-constrained simulation pipeline. `PROFILE`, `ATTACK`, and `OPINION` are encoded as hierarchical ontologies; each execution instantiates a factorial attack-opinion surface over a deterministic panel of synthetic profiles, elicits baseline and post-exposure opinions, and then estimates directional adversarial effectivity rather than undirected change magnitude. The broader program goal is not to claim that large language model agents are substitutes for human subjects, but to produce a transparent, auditable measurement object that can later be compared against human data using criterion and construct validity tests.

The present study is a deliberate stress test of that analysis stack. It focuses on a set of representative ontology leaves — one attack per mechanism family and a cross-domain opinion sample — and introduces a more advanced conditional susceptibility dashboard and a richer network analysis module aimed at distinguishing dense local hubs from bridge variables in the mixed continuous-plus-dummy profile panel. The paper therefore asks a narrower question than the long-run research program: *does the upgraded analysis pipeline respond defensibly when the execution substrate is degraded, and what interpretable structure can still be recovered from deterministic components of the design?*

Methods

Study Design

The present study used an attacked-only full factorial design with 25 deterministic pseudoprofiles, four attack vectors, and ten opinion leaves ($25 \times 4 \times 10 = 1{,}000$ scenarios). The four attacks were one leaf from each mechanism family: Misleading Narrative Framing, Astroturf Comment Wave, Fear Appeal Scapegoating Post, and LLM Chatbot Personalized Persuasion. The ten opinion leaves were sampled across defense and national security, foreign policy and geopolitics, immigration and citizenship, and information integrity and platforms.

The active profile panel contained 91 features spanning Big Five traits and facets, socioeconomic standing, employment, institutional trust, ideological positioning, political engagement, online behavior, social capital, age, sex, education, and news-diet/platform indicators. Earlier pilot-only inventories that weakened interpretability and regularized fit stability were excluded from the active feature registry. Profiles were generated deterministically, so the feature matrix and all downstream network results are exactly reproducible from the configuration and random seed.

Outcome Definition and Model Stack

For each attack-opinion-profile triple, the primary outcome was adversarial effectivity:

$$\mathrm{AE}_{iak} = (\mathrm{post}_{iak} - \mathrm{baseline}_{iak}) \times d_k, \tag{1}$$

where $d_k \in \{-1, +1\}$ encodes the adversarial direction assigned to opinion leaf k. Positive values indicate movement in the adversarially intended direction; negative values indicate resistance or backfire.

The analysis stack comprised five layers. First, profile-level variance partitioning used ICC(1) to estimate the share of variance attributable to stable between-profile differences (Shrout and Fleiss, 1979). Second, a repeated-outcome SEM regressed attack-averaged effectivity indicators on profile moderators, treated as a small-panel descriptive diagnostic rather than a confirmatory latent model (Rosseel, 2012). Third, task-specific ridge models were fit for each attack-opinion cell, then aggregated into a conditional susceptibility score using reliability weights derived from sample size and cross-validated error. Fourth, bootstrap resampling at the profile level quantified rank instability and coefficient uncertainty; false-discovery-rate control used the Benjamini-Hochberg procedure (Benjamini and Hochberg, 1995). Fifth, a profile feature network was computed from the deterministic profile matrix using Spearman correlations with threshold $|\rho| \geq .15$.

Advanced Profile Feature Network

The network analysis is designed for a mixed hierarchical panel in which continuous trait scales coexist with one-hot categorical indicators. Edge weights use $|\rho|$, while path-based metrics use distance $= 1/|\rho|$. Alongside degree, eigenvector centrality, PageRank, closeness, and betweenness, the network export includes weighted degree (strength), participation coefficient, bridge ratio, within-module z score, sign balance, positive-versus-negative correlation strength, community size, and k-core index. In the interactive explorer, community convex hulls demarcate detected clusters, force-directed layout enables manual spreading of dense regions, and lasso multi-selection supports cohort-level metric comparison. This distinction matters methodologically: in a dense mixed panel, high degree alone does not indicate whether a variable is a local hub within one ontology block or a bridge that ties otherwise separated families together.

Execution-Integrity Diagnostics

Execution quality was audited explicitly. Baseline and post-exposure stages record fallback usage, realism and plausibility review outcomes, and downstream robustness sum-

maries. This is central to the inferential posture of the platform: if execution fails, the analysis should reveal that failure and downgrade claims instead of allowing visually polished but methodologically indefensible conclusions to pass unchallenged. That design principle is especially important in adversarial-simulation work, where reviewers are rightly skeptical of overfit dashboards and overly certain behavioral narratives.

Results

Execution Integrity Dominates Interpretation

The present study completed structurally but failed substantively at the opinion-generation layer. Baseline and post-exposure assessment used deterministic fallback outputs for every row, and attack-review scores remained at zero throughout the realism pipeline. Those diagnostics alone imply that the present study cannot support substantive claims about persuasion dynamics. The appropriate interpretation is therefore constrained from the start: this is a validation of the upgraded analysis stack under degraded execution, not a behavioral finding about real or human-like susceptibility.

The descriptive outcomes are consistent with that interpretation. Mean absolute displacement was 34.23 points, mean adversarial effectivity was 0.10, and 49.6% of rows showed positive adversarial effectivity. The pipeline produced non-trivial movement magnitudes but essentially random directional success. Had the pipeline reported strong profile moderators under those conditions, that would have been a methodological red flag.

Figure 1: Study factorial design. The study crossed 25 deterministic pseudoprofiles with four attack leaves and ten opinion leaves drawn across four policy domains. All observations are exposed profile-attack-opinion triples.

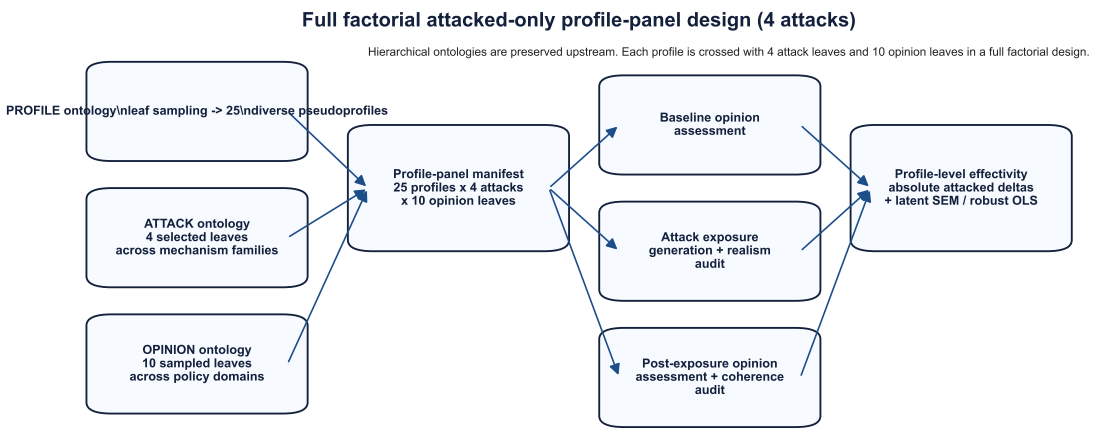

Predictive Diagnostics Correctly Collapse

Variance partitioning and out-of-sample diagnostics all converged on the same conclusion. ICC(1) for absolute shift was 0.011, indicating that almost all variance in $|\Delta|$ sat within profiles across tasks rather than between profiles. Ridge cross-validated R^2 was -0.175, random-forest out-of-bag R^2 was -0.194, and the elastic net selected zero features.

Together these results show that the upgraded moderation stack is behaving defensibly: under fallback-generated opinion data, it refuses to manufacture predictive structure.

This negative result is useful. It demonstrates that the methodological guardrails do what they are supposed to do under adverse conditions. A less defensible system might still return a visually compelling ranking, a handful of unstable coefficients, and an overconfident narrative. Here, the diagnostics instead converge on the absence of reliable profile-level signal.

Figure 2: Opinion displacement and adversarial effectivity distribution. Each point represents one attacked scenario ($n = 1{,}000$). Colour encodes attack vector. Opinion displacement was large across all attack types and opinion leaves, while adversarial effectivity was near zero and directionally random — consistent with fully fallback-generated opinion data.

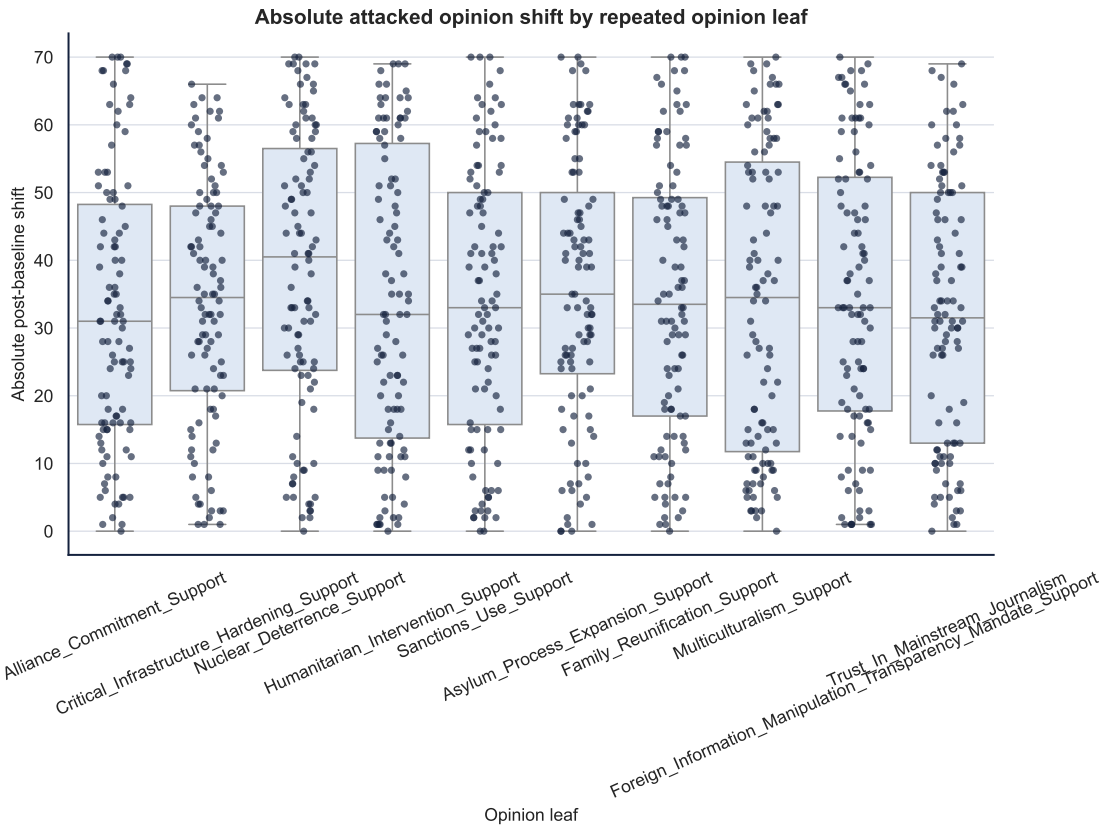

Network Structure Remains Interpretable

Unlike the opinion-generation stages, the profile feature network is unaffected by API failure because it is computed from the deterministic profile matrix. The graph contained 91 nodes and 1,807 edges at $|\rho| \geq .15$, with density = .441, five communities, and modularity = .313. That topology indicates a moderately clustered but still highly cross-connected feature space. The most connected regions concentrated in online-behavior indicators, institutional-trust variables, and Big Five openness/extraversion blocks, which is consistent with the way the deterministic profile generator reuses coherent latent structure across related features.

The more informative result is not simply that the graph is dense, but that bridge diagnostics separate local hubs from cross-family connectors. Variables with high partici-

pation or bridge ratio are not merely central; they tie otherwise distinct ontology families together. In practice, these are the nodes worth scrutinizing when reviewing redundancy, multicollinearity, or hidden coupling between dummy-coded demographics and continuous psychographic scales. This is exactly the kind of diagnostic that a reviewer would ask for in a mixed hierarchical feature panel.

Figure 3: Profile-moderator coefficient forest. OLS coefficient estimates on profile-mean adversarial effectivity with bootstrap 95% confidence intervals ($B = 200$). Under fallback-generated opinion data, all intervals cross zero — consistent with the absence of reliable profile-level signal. This diagnostic is the correct output from the pipeline under degraded execution.

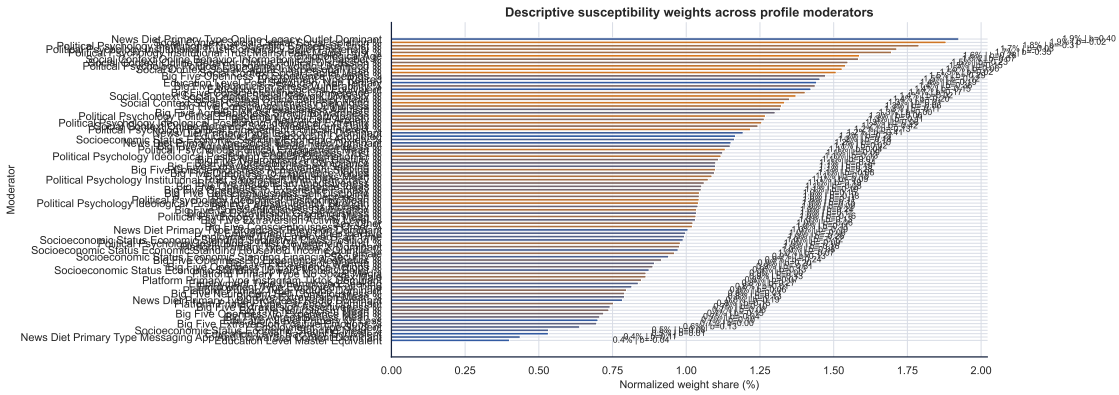

Why the Negative Result Is Analytically Important

Peer review should be harsher on synthetic susceptibility studies than on ordinary exploratory dashboards, not more permissive. The present study therefore passes only if the upgraded system remains conservative when the data-generation chain degrades. On that criterion, the study is informative: the inferential layers that depended on modelled opinion responses degraded toward noise, while the deterministic layer that depended only on feature construction remained stable and interpretable. That separation between *execution-sensitive* and *execution-invariant* components is itself a useful validation outcome.

Discussion

The main contribution of the present study is that it demonstrates analytical self-restraint. The dashboard, tables, and manuscript no longer treat every completed execution as if it merited substantive inference. Instead, the platform now exposes execution quality, shows when predictive fit collapses, and routes attention toward the parts of the pipeline that remain interpretable. That is the correct methodological stance for an ontology-driven simulation framework that will eventually be compared against real opinion data.

Three reviewer-facing points follow directly. First, this study does *not* justify any claim that certain personalities or demographic groups are more susceptible to adversarial persuasion. The fallback-dominated opinion stages invalidate that inference. Second, the network analysis should not be mistaken for evidence about persuasion outcomes. It characterizes the geometry of the modeled feature panel, not the success of attacks. Third, the upgraded network explorer is substantively useful because it diagnoses redundancy, bridge structure,

and mixed-feature coupling before any higher-stakes modeling claims are made.

This perspective is also operationally practical. A replication with funded API access from stage 02 onward can now reuse the upgraded stack without further architectural changes. If that replication still yields near-zero predictive fit, the correct conclusion will be that the broader multi-domain attack surface contains weak profile-level signal. If instead real opinion-generation outputs recover stable task structure or bridge-sensitive moderators, the upgraded diagnostics and dashboard are already in place to show that signal defensibly.

Conclusion

The present study should be read as a validation of analytical infrastructure, not of substantive persuasion findings. The study confirms that the upgraded analysis stack can:

1. preserve a reproducible ontology-driven factorial design over profiles, attacks, and opinions;
2. surface execution failure rather than masking it with overconfident moderator claims;
3. quantify predictive collapse through ICC, ridge, random-forest, elastic-net, and bootstrap diagnostics; and
4. recover interpretable deterministic structure through an advanced profile feature network with hub-versus-bridge diagnostics and an interactive explorer that supports force-directed layout, community hull visualization, lasso multi-selection, and shortest-path analysis.

That is a stronger result than a noisy list of nominally significant coefficients would have been. The next scientifically meaningful step is straightforward: restore API access, re-execute stages 02–09 for the same design, and then evaluate whether the upgraded diagnostics support any profile-level susceptibility claims under normal execution.

Dual-Use Ethics Statement

This project models mechanisms that could, in principle, be repurposed for adversarial targeting. The intended use is defensive: auditing vulnerability structure, stress-testing analytical claims, and identifying where resilience interventions may be most needed. No real individuals are modeled; all profiles are synthetic; and the present study yields no operationally actionable targeting signal. Any future deployment on real data would require ethics review, access controls, and explicit prohibition of offensive use.

Reproducibility Statement

The full simulation pipeline is implemented in Python and driven by hierarchical JSON ontologies. All analyses are reproducible from the ontology files and the random seed (seed

= 100, bootstrap samples = 200). Stage-level artifacts — delta summaries, ICC outputs, OLS coefficient tables, SEM fit indices, bootstrap CSVs, ridge reliability weights, and hierarchical importance breakdowns — are stored under the run's `stage_outputs/` directory with per-stage `manifest.json` files for hash-level verification. The deterministic profile feature network is reproducible independent of API availability. The opinion-generation components used fallback outputs in the present study; to obtain real opinion data, restore API credits and re-execute from stage 02.

Supplementary Materials

Figure 4: Supplementary Figure S1. Baseline versus post-attack opinion scores. Each point represents one attacked scenario; colour encodes adversarial effectivity. The diagonal concentration reflects the dominance of fallback-generated scores.

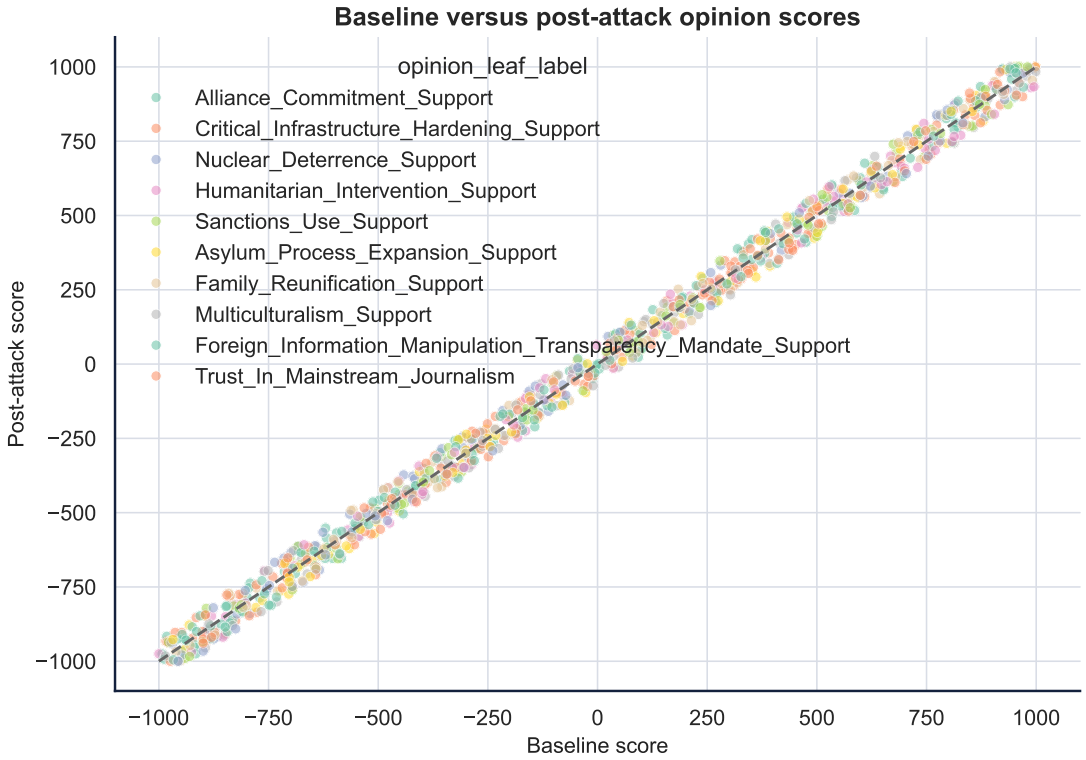

Figure 5: Supplementary Figure S2. Per-profile attacked-effectivity heatmap. Rows are profiles (ordered by susceptibility index); columns are attack-opinion task combinations.

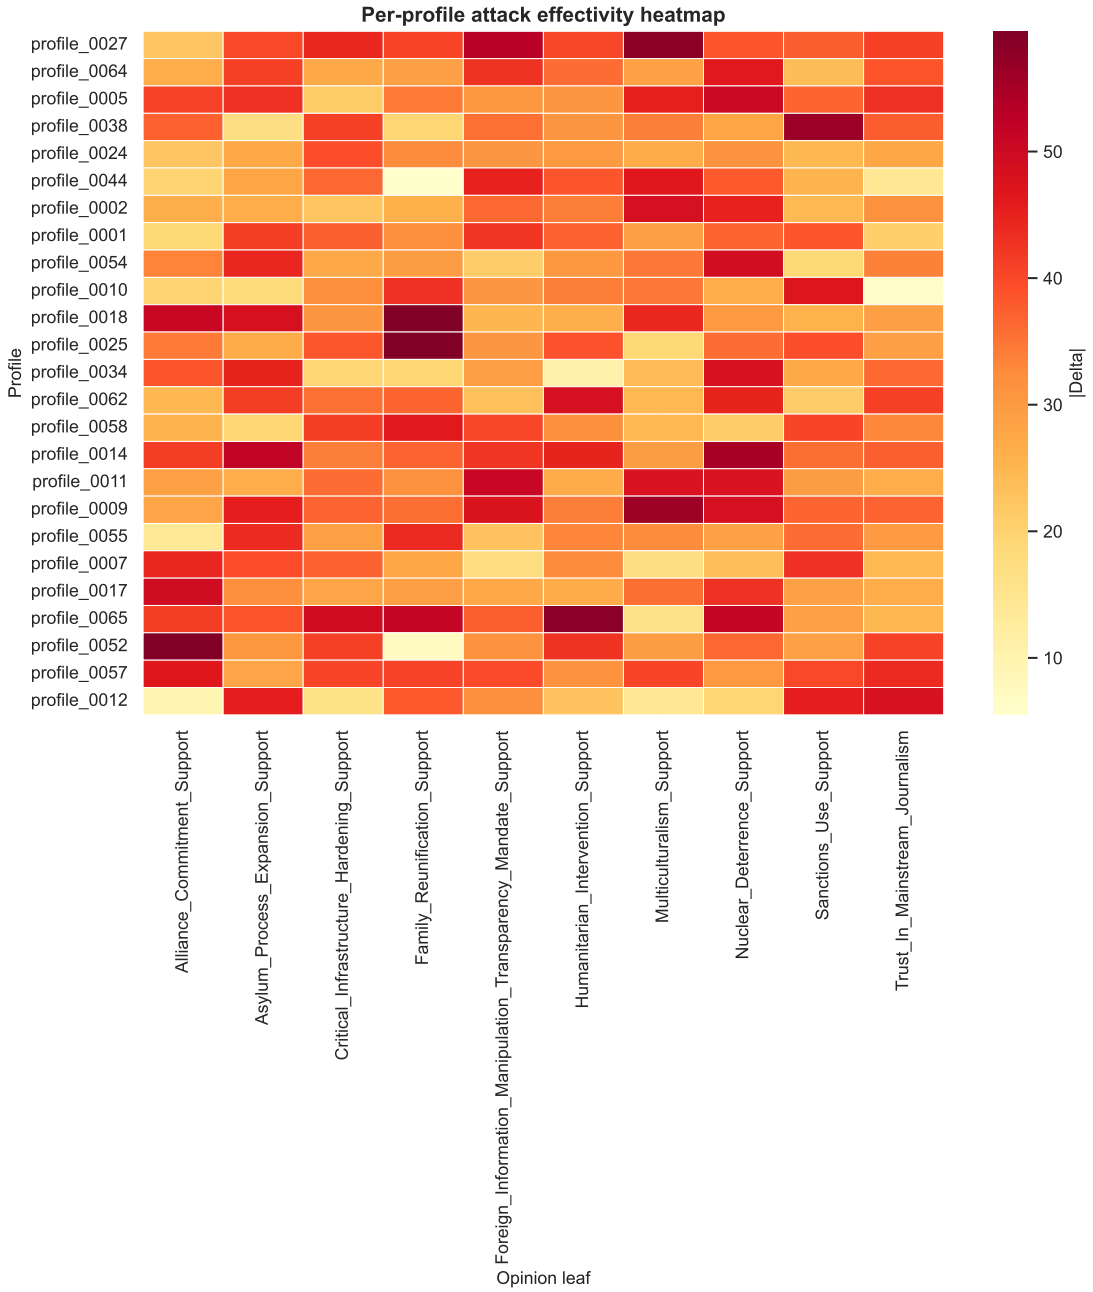

Figure 6: Supplementary Figure S3. Conditional susceptibility distribution across profiles with bootstrap confidence intervals for the top and bottom quintile.

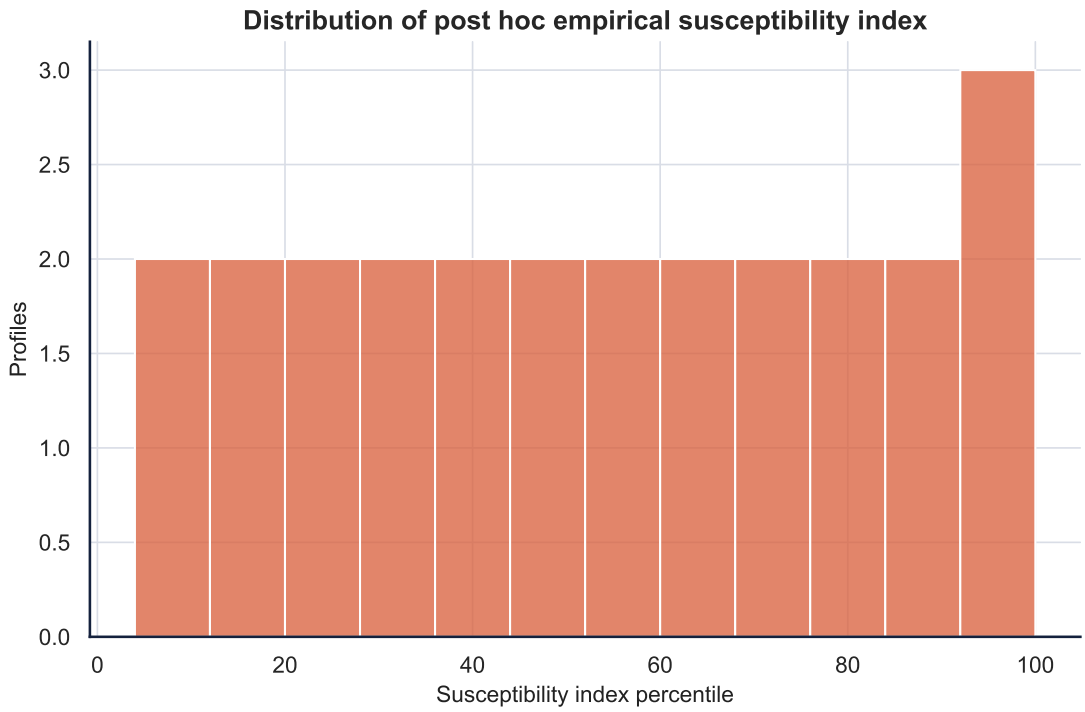

Figure 7: Supplementary Figure S4. Attack-comparison panel (adversarial effectivity by attack vector). No attack type produces systematically positive mean adversarial effectivity under fallback execution.

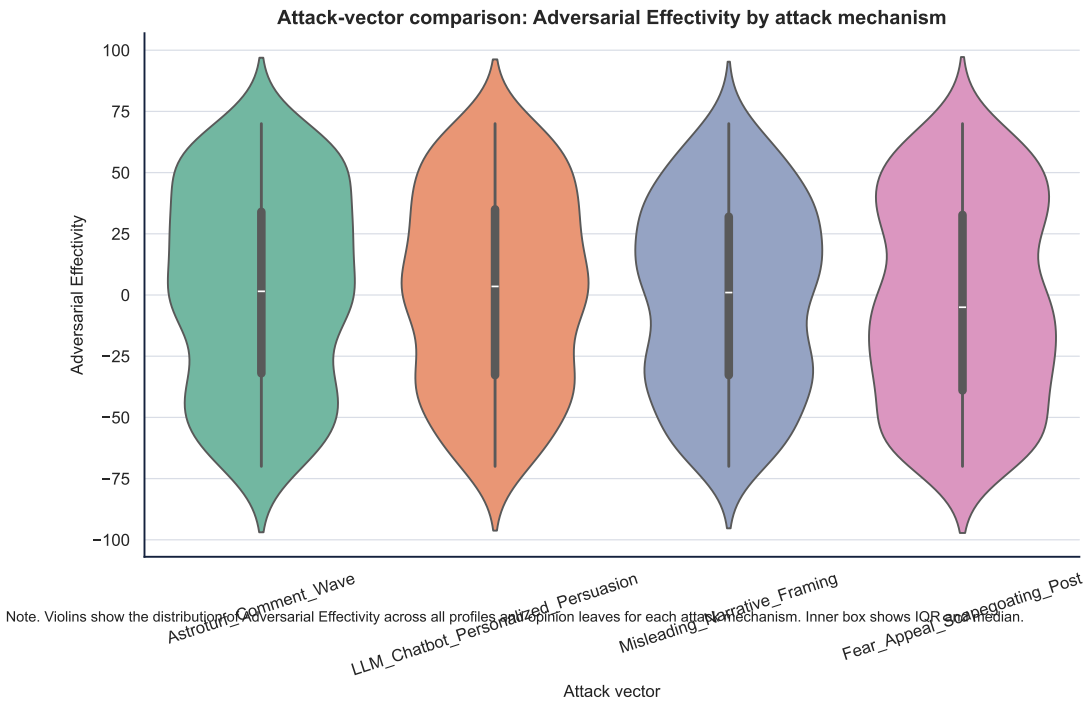

Acknowledgements

Table 1: Ontology leaves used in run_10.

Ontology	Role	Leaf
ATTACK	Factorial attack leaf	ATTACK_VECTORS > Social_Media_Misinformation > Misleading_Narrative_Framing
ATTACK	Factorial attack leaf	ATTACK_VECTORS > Coordinated_Amplification > Astroturf_Comment_Wave
ATTACK	Factorial attack leaf	ATTACK_VECTORS > Manipulative_Persuasion > Fear_Appeal_Scapegoating_Post
ATTACK	Factorial attack leaf	ATTACK_VECTORS > Synthetic_Media_Generation > LLM_Chatbot_Personalized_Persuasion
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Defense_and_National_Security > Alliance_Commitment_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Defense_and_National_Security > Critical_Infrastructure_Hardening_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Defense_and_National_Security > Nuclear_Deterrence_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Foreign_Policy_and_Geopolitics > Humanitarian_Intervention_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Foreign_Policy_and_Geopolitics > Sanctions_Use_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Immigration_and_Citizenship > Asylum_Process_Expansion_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Immigration_and_Citizenship > Family_Reunification_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Immigration_and_Citizenship > Multiculturalism_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Information_Integrity_and_Platforms > Foreign_Information_Manipulation_Transparency_Mandate_Support
OPINION	Repeated indicator leaf	Issue_Position_Taxonomy > Information_Integrity_and_Platforms > Trust_In_Mainstream_Journalism

Note. Only ontology leaves are sampled for estimation. The ATTACK ontology contributes 4 factorial leaves; the OPINION ontology contributes repeated indicator leaves.

Table 2: Core and exploratory profile moderator comparison.

Moderator	Role	Multivariate b	Multivariate p	Univariate b	Univariate p	Ontology group	Normalized weight %
Big Five Openness To Experience Mean %	core	-2.142	0.122	-1.834	0.028	Big Five: Openness	0.720
Big Five Neuroticism Mean %	core	1.604	0.190	1.076	0.162	Big Five: Neuroticism	0.738
Sex Other	core	5.037	0.239	1.748	0.425	Demographics: Sex	1.023
Big Five Conscientiousness Mean %	core	1.264	0.427	0.494	0.689	Big Five: Conscientiousness	1.087
Sex Female	core	3.043	0.494	0.454	0.844	Demographics: Sex	0.961
Big Five Extraversion Mean %	exploratory	0.650	0.673	-0.248	0.829	Big Five: Extraversion	0.791
Big Five Agreeableness Mean %	exploratory	0.543	0.719	-0.307	0.786	Big Five: Agreeableness	0.707

Note. Core terms are entered into the latent SEM or primary multivariate profile model. Normalized weight percentages summarize each moderator's share of the total fitted importance after accounting for the moderator's observed variability.

Table 3: Assumption register and peer-review risk register.

Type	Item	S
Assumption	Pilot sample size is sufficient for stable moderation estimates.	c
Assumption	Attack-only design is intentional and aligned with the moderation question.	c
Assumption	Opinion measurement has enough numeric resolution.	c
Assumption	Attack text realism is adequate.	r
Assumption	SEM converges with acceptable fit.	c
Assumption	LLM outputs are reproducible and auditable.	c
Critique	Synthetic LLM agents may not represent human cognition or causal response behavior.	a
Critique	Adversarial content may be unrealistic or too generic.	a
Critique	Model dependence on a single LLM could bias outputs.	a
Critique	Insufficient sample size inflates uncertainty and may destabilize SEM.	a
Critique	Attack-only designs cannot estimate the incremental effect of exposure relative to a no-attack counterfactual.	a

Note. The pilot is methodological and exploratory. Risks are surfaced explicitly so later scale-up runs can target the current design bottlenecks.

The authors thank the Ghent University Department of Applied Mathematics, Computer Science and Statistics for institutional support. Computation used Python 3.12 with NumPy, pandas, scikit-learn, statsmodels, semopy, Plotly, NetworkX, and seaborn.

References

Benjamini, Y. and Hochberg, Y. (1995). Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1):289–300.

Bennett, W. L. and Livingston, S. (2018). The disinformation order: Disruptive communication and the decline of democratic institutions. *European Journal of Communication*, 33(2):122–139.

Kozyreva, A., Lewandowsky, S., and Hertwig, R. (2020). Citizens versus the internet: Confronting digital challenges with cognitive tools. *Psychological Science in the Public Interest*, 21(3):103–156.

Lazer, D. M. J., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., Metzger, M. J., Nyhan, B., Pennycook, G., Rothschild, D., Schudson, M., Sloman, S. A., Sunstein, C. R., Thorson, E. A., Watts, D. J., and Zittrain, J. L. (2018). The science of fake news. *Science*, 359(6380):1094–1096.

Lewandowsky, S., Ecker, U. K. H., and Cook, J. (2017). Beyond misinformation: Understanding and coping with the post-truth era. *Journal of Applied Research in Memory and Cognition*, 6(4):353–369.

Matz, S. C., Kosinski, M., Nave, G., and Stillwell, D. J. (2017). Psychological targeting as an effective approach to digital mass persuasion. *Proceedings of the National Academy of Sciences*, 114(48):12714–12719.

Table 4: Reproducibility manifest for run_10.

Field	Value
run_id	run_10
paper_title	PILOT: Inter-individual Differences in Susceptibility to Cyber-manipulation: A Multi-agent Simulation Approach with High-dimension State Space of Political Opinions
output_root	/Users/stijnvanseveren/PythonProjects/Paper_CaseStudiesAnalysisExperimentalData/evaluation
report_root	/Users/stijnvanseveren/PythonProjects/Paper_CaseStudiesAnalysisExperimentalData/research_r
report_assets_root	/Users/stijnvanseveren/PythonProjects/Paper_CaseStudiesAnalysisExperimentalData/research_r
n_scenarios	10
n_profiles	25
seed	100
attack_ratio	1.000
attack_leaf	ATTACK_VECTORS > Social_Media_Misinformation > Misleading_Narrative_Framing
attack_leaves	Misleading_Narrative_Framing,Astroturf_Comment_Wave,Fear_Appeal_Scapegoating_Post,LL
focus_opinion_domain	
max_opinion_leaves	10
profile_candidate_multiplier	3
primary_moderator	posthoc_profile_susceptibility_index
bootstrap_samples	200
use_test_ontology	True
ontology_root	/Users/stijnvanseveren/PythonProjects/Paper_CaseStudiesAnalysisExperimentalData/src/backen
openrouter_model	mistralai/mistral-small-3.2-24b-instruct
temperature	0.150
max_repair_iter	3
profile_generation_mode	deterministic
self_supervise_attack_realism	True
realism_threshold	0.700
self_supervise_opinion_coherence	True
coherence_threshold	0.700
generate_visuals	True
export_static_figures	True
build_report	True
resume_from_stage	01
stop_after_stage	09
save_raw_llm	True
timeout_sec	90
max_concurrency	20
run_stage_checks	False
stage_check_scenarios	3

Note. The manifest captures the full pipeline configuration used to generate the pilot outputs and manuscript assets.

Table 5: Leaf-specific path-SEM coefficients.

Outcome leaf	Moderator
Adversarial Delta Indicator Alliance Commitment Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Alliance Commitment Support	Big Five Openness To Experi
Adversarial Delta Indicator Alliance Commitment Support	Big Five Conscientiousness M
Adversarial Delta Indicator Alliance Commitment Support	Sex Female
Adversarial Delta Indicator Alliance Commitment Support	Sex Other
Adversarial Delta Indicator Alliance Commitment Support	Mean Baseline Abs Score
Adversarial Delta Indicator Asylum Process Expansion Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Asylum Process Expansion Support	Big Five Openness To Experi
Adversarial Delta Indicator Asylum Process Expansion Support	Big Five Conscientiousness M
Adversarial Delta Indicator Asylum Process Expansion Support	Sex Female
Adversarial Delta Indicator Asylum Process Expansion Support	Sex Other
Adversarial Delta Indicator Asylum Process Expansion Support	Mean Baseline Abs Score
Adversarial Delta Indicator Critical Infrastructure Hardening Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Critical Infrastructure Hardening Support	Big Five Openness To Experi
Adversarial Delta Indicator Critical Infrastructure Hardening Support	Big Five Conscientiousness M
Adversarial Delta Indicator Critical Infrastructure Hardening Support	Sex Female
Adversarial Delta Indicator Critical Infrastructure Hardening Support	Sex Other
Adversarial Delta Indicator Critical Infrastructure Hardening Support	Mean Baseline Abs Score
Adversarial Delta Indicator Family Reunification Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Family Reunification Support	Big Five Openness To Experi
Adversarial Delta Indicator Family Reunification Support	Big Five Conscientiousness M
Adversarial Delta Indicator Family Reunification Support	Sex Female
Adversarial Delta Indicator Family Reunification Support	Sex Other
Adversarial Delta Indicator Family Reunification Support	Mean Baseline Abs Score
Adversarial Delta Indicator Foreign Information Manipulation Transparency Mandate Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Foreign Information Manipulation Transparency Mandate Support	Big Five Openness To Experi
Adversarial Delta Indicator Foreign Information Manipulation Transparency Mandate Support	Big Five Conscientiousness M
Adversarial Delta Indicator Foreign Information Manipulation Transparency Mandate Support	Sex Female
Adversarial Delta Indicator Foreign Information Manipulation Transparency Mandate Support	Sex Other
Adversarial Delta Indicator Foreign Information Manipulation Transparency Mandate Support	Mean Baseline Abs Score
Adversarial Delta Indicator Humanitarian Intervention Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Humanitarian Intervention Support	Big Five Openness To Experi
Adversarial Delta Indicator Humanitarian Intervention Support	Big Five Conscientiousness M
Adversarial Delta Indicator Humanitarian Intervention Support	Sex Female
Adversarial Delta Indicator Humanitarian Intervention Support	Sex Other
Adversarial Delta Indicator Humanitarian Intervention Support	Mean Baseline Abs Score
Adversarial Delta Indicator Multiculturalism Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Multiculturalism Support	Big Five Openness To Experi
Adversarial Delta Indicator Multiculturalism Support	Big Five Conscientiousness M
Adversarial Delta Indicator Multiculturalism Support	Sex Female
Adversarial Delta Indicator Multiculturalism Support	Sex Other
Adversarial Delta Indicator Multiculturalism Support	Mean Baseline Abs Score
Adversarial Delta Indicator Nuclear Deterrence Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Nuclear Deterrence Support	Big Five Openness To Experi
Adversarial Delta Indicator Nuclear Deterrence Support	Big Five Conscientiousness M
Adversarial Delta Indicator Nuclear Deterrence Support	Sex Female
Adversarial Delta Indicator Nuclear Deterrence Support	Sex Other
Adversarial Delta Indicator Nuclear Deterrence Support	Mean Baseline Abs Score
Adversarial Delta Indicator Sanctions Use Support	Big Five Neuroticism Mean %
Adversarial Delta Indicator Sanctions Use Support	Big Five Openness To Experi
Adversarial Delta Indicator Sanctions Use Support	Big Five Conscientiousness M
Adversarial Delta Indicator Sanctions Use Support	Sex Female
Adversarial Delta Indicator Sanctions Use Support	Sex Other
Adversarial Delta Indicator Sanctions Use Support	Mean Baseline Abs Score
Adversarial Delta Indicator Trust In Mainstream Journalism	Big Five Neuroticism Mean %
Adversarial Delta Indicator Trust In Mainstream Journalism	Big Five Openness To Experi
Adversarial Delta Indicator Trust In Mainstream Journalism	Big Five Conscientiousness M
Adversarial Delta Indicator Trust In Mainstream Journalism	Sex Female
Adversarial Delta Indicator Trust In Mainstream Journalism	Sex Other
Adversarial Delta Indicator Trust In Mainstream Journalism	Mean Baseline Abs Score

Pennycook, G. and Rand, D. G. (2021). The psychology of fake news. *Trends in Cognitive Sciences*, 25(5):388–402.

Roozenbeek, J., van der Linden, S., Goldberg, B., Rathje, S., and Lewandowsky, S. (2022). Psychological inoculation improves resilience against misinformation on social media. *Science Advances*, 8(34):eabo6254.

Rosseel, Y. (2012). lavaan: An R package for structural equation modeling. *Journal of Statistical Software*, 48(2):1–36.

Shrout, P. E. and Fleiss, J. L. (1979). Intraclass correlations: Uses in assessing rater reliability. *Psychological Bulletin*, 86(2):420–428.